Understanding Diverging Perceptions of Artificial Intelligence (AI): The Role of the Welfare State

Stellar Zhang*

2025-12-17

Abstract

This paper explores determinants of the perceived labour market impacts of artificial intelligence. I develop a theory connecting the generosity of the national welfare state with whether these impacts are perceived as economic threats or opportunities. I argue that when an individual will be protected by the welfare state in the event of a major economic shock, they are more likely to view AI in light of its potential opportunities for increased productivity and performance. In comparison, an individual who would not be protected from dislocative shocks is more likely to view AI as a threat, creating unemployment or inequality. Using data from the 2024 wave of the OECD Risks that Matter survey, I find moderate evidence in support of my theory, but only through subjective, rather than objective, assessments of the welfare state's generosity.

*Department of Political Science, McGill University. Email: stellar.zhang@mail.mcgill.ca.

1. Introduction

The rapid pace of artificial intelligence (AI)'s evolving capabilities has sparked both fear and optimism about its potential to transform labour markets.[1] Some estimates show that two-thirds of occupations in the US and Europe are exposed to AI to at least some extent, and that one-fourth of current work could be substituted for AI entirely (Briggs and Kondani, 2023). Other studies find that with the use of AI, over half (56%) of all tasks could be completed faster without sacrificing quality (Elondou et al., 2024). While these are still only forecasts, AI adoption is already underway in a substantial portion of the economy, with one estimate putting it at 18% of the employment-adjusted economy (McElheran et al., 2024). As the technology improves, it is likely that these numbers will only increase.

Despite the scale of these estimates, the exact nature of AI's economic impact remains highly uncertain. More pessimistic views argue that the widespread adoption of AI will cause large-scale economic displacement and unemployment as it outpaces human labour in terms of cost, efficiency, or performance. Optimists argue that, like previous technological transformations, AI could bolster productivity, innovation, and ultimately prosperity. This debate is vigorous and still ongoing, much of it drawing upon previous technological transformations such as automation or digitialisation to inform future predictions.

However, unlike previous shocks whose effects were concentrated in certain industries, skill groups, or countries, the effect of AI is much more wide-ranging. The broad consensus is that most occupations exhibit at least some degree of exposure to varying degrees (Elondou et al. 2023, Briggs and Kondani, 2023). However, occupational exposure has a dual meaning. Under a task-based model of occupations, AI has the potential to complement human workers, leading to gains in efficiency, productivity, and performance, but it can also risk substituting human labour altogether (Autor et al., 2003). Those whose tasks largely overlap with current or future AI capabilities will experience displacement or other adverse outcomes, while those whose work would benefit from it would enjoy higher efficiency. This makes AI either a threat or opportunity to an occupation. Uniquely, these are not mutually exclusive categorizations. The same occupation can, and indeed often does, face both high complementarity potential and high substitution risk, such as professional and managerial positions (Pizzinelli et al., 2023).

A large body of literature has documented how the previous economic transformations, such as trade liberalization, digitialization, or labour-saving automation, have resulted in significant distributional consequences, creating clear winners and losers (Autor et al, 2013; Acemoglu and Restrepo, 2019). Previous advances in technology favoured educated and skilled workers over the unskilled, particularly those with a college degree (Katz and Murphy, 1992; Gallego et al., 2022). The globalization and import competition enabled by the ICT revolution harmed low-skilled workers in capital-rich countries, whose work in industries such as manufacturing were moved offshore (Antràs, 2020).

[1] For the purposes of this paper, I follow Briggs and Kondani (2023) in defining the current technologies broadly categorized as AI as those possessing the following features: (1) general, as opposed to narrow, use cases; (2) the ability to create new output rather than simply describe existing information; and (3) demonstrating an understanding of and ability to convey information in natural human language, as well as images or video.

Transformations such as these have disenfranchised workers whose jobs are relocated, automated, or even rendered entirely obsolete. Many studies then link this disenfranchisement to negative political consequences, primarily populist backlashes (Autor et al. 2020, Baccini and Weymouth, 2021; Colantone and Stanig, 2018; Rodrik, 2021). The longstanding embedded liberalism hypothesis posits that to avoid these consequences, governments must compensate citizens for such market shocks through welfare (Ruggie, 1982). Such compensatory policies include, but are not limited to, the provision of unemployment benefits, vocational retraining schemes, or even redistribution.

These policies aim to mitigate the sense of risk among the vulnerable, and the evidence shows that they can be moderately successful (Vlandas and Halikiopoulou, 2021). However, much of this work examines retrospective effects: how the effects of past economic disruptions must be retroactively offset. Less is known about the anticipatory dimension: how citizens' existing institutional environment shapes their perception of future risks, even before those shocks have fully realized.

My paper builds on a body of literature examining predictors of perceived AI risk. Recent work has focused on demographic predictors such as education, gender, and employment sector, and their effect on AI outlooks; or else the effect of informational treatments on AI-related policy preferences (Borwein et al, 2025; Zhang, 2022). However, there remains little evidence on how national, institutional contexts mediate these perceptions. This paper's contribution will be to examine the welfare state conditions citizens' perceptions of AI as either a threat or an opportunity, and, by extension, their policy preferences, using large-N survey data from the OECD.

I find mixed evidence in support of my theory. I find statistically significant results when examining an individual's subjective level of welfare access, i.e. when asked to hypothetically evaluate how likely it would be that the government would support them in the event of unemployment. These effectively predict their outlooks towards AI in the expected directions. However, these do not map precisely onto the objective level of government spending, as captured by the percentage of GDP spent on welfare, nor do they predict policy preferences.

The rest of this paper proceeds as follows. Section 2 lays out the motivations and theory behind my main argument, as well as my expected findings. Section 3 describes the data and empirical methods used to test the argument. Section 4 discusses my findings, revisits my initial expectations, and concludes.

2. Theoretical Underpinnings

I develop a theory linking the welfare state to diverging AI perceptions. My argument relies on three key assumptions:

1. Workers understand, at least to some extent, their unique levels of economic exposure to AI.
2. Workers turn to the state and expect protection from shocks by providing welfare, unemployment benefits, or retraining schemes.

3. Workers then account for the state's role in welfare provision when evaluating their future predictions about the impact of AI, and their subsequent preferences.

As described above, AI can present simultaneously as a threat or as an opportunity. It can automate menial or repetitive work, raising productivity and enabling people to have more free time. It could also create jobs yet unheard of. Previous major transformations, such as the ICT revolution, created new technological needs, resulting in jobs for software engineers, cybersecurity analysts, and patent lawyers, among many others (Autor 2024). New jobs like these create both positive productivity effects as well as new demand for human labour (Acemoglu and Restropo, 2019).

However, AI can also pose an economic threat. It is best at performing routine cognitive tasks, such as writing computer code or searching through databases for information (Webb, 2020). However, it cannot currently perform manual or emotional tasks, such as caregiving, nor high-level tasks such as strategic decision-making. The widespread adoption of AI risks hollowing out the class of knowledge workers who typically perform mid-level cognitive work. At the same time, its benefits are concentrated among the highest earners (Pizzinelli et al., 2023). This would result higher wage polarization and an increase in income inequality, a dynamic that has historically played out with previous economic disruptions (Autor 2015; Michaels et al. 2014).

Unsure about the evolving impact of AI, I posit that workers turn to the state for protection. Those who believe that the state adequately protects their social welfare and economic security have less to fear about the dislocative shocks of AI, resulting in reduced threat perceptions and higher opportunity perceptions. These workers feel sufficiently insulated: should they lose their job to AI, they know they will be supported by the government in retraining, finding a new job, or receiving supplemental income in the interim. Their level of social protection therefore enables them to be optimistic about the opportunity AI presents. On the contrary, workers who believe that the state is inadequate in protecting them are expected to more strongly perceive AI as a threat and less as an opportunity.

I argue that this mechanism potentially operates at two different levels. The first is the objective level of welfare support provided by a state, operationalised as the percentage of GDP spent on welfare benefits. The second is the subjective individual level, measured as an individual's perceived level of access to those benefits. Although these two are likely to be correlated, with individuals in more generous welfare states broadly expressing more approving opinions, they are not entirely the same. The final effect I assess is how one's perceived vulnerability to AI magnifies the existing relationship. Given this, I hypothesize the following, which I term the divergence theory of AI perceptions.

H1. Workers in more generous welfare states (defined as those that spend a greater portion of GDP on social support) are less likely to see AI as a threat and more likely to see it as an opportunity, compared to those in less generous welfare states.

H2. Workers who believe they can access unemployment benefits from the state are less likely to see AI as a threat and more likely to see it as an opportunity, compared to those who believe they cannot.

H3. The diverging relationships in H1 and H2 will be especially strong among workers who self-identify as being at high risk of being replaced by AI in the foreseeable future.

I then argue that these expectations inform policy preferences. Governments have a range of options when facing the introduction of AI into their economies. Regulation would entail restriction on the development of AI, including introducing a tax on technology companies (a so-called robot tax) or placing tariffs on imports of foreign technology. Comparatively, integration would include policies aimed to better incorporate the role of AI in the economy, such as investing into digital infrastructure (such as broadband networks or data processing centres), providing subsidies to the hardest-hit firms, or retraining workers. In line with diverging perceptions, I argue that the welfare state creates diverging policy preferences as well.

H4: Workers in generous welfare states express stronger preferences for AI integration and weaker preferences for regulation.

H5: Workers who believe they can access unemployment benefits express stronger preferences for AI integration and weaker preferences for regulation.

I test these hypotheses with large-N survey evidence from OECD countries in the following sections.

3. Data and Methodology

The data come from the 2024 wave of the OECD's *Risks that Matter (RTM)* survey. It asks over 27,000 respondents across the 27 member countries about their perceived social and economic risks, as well as preferences for government social policies. Research has found that the labour market impact of AI in advanced economies, which constitute the majority of those in the OECD, is more polarized than in emerging economies, making it an ideal case to study the divergence of opinion (Pizzinelli et al., 2023).

Respondents are also asked about their preferences towards government policies concerning these risks, including social welfare, foreign policy, or market interventions. Individual-level demographic data is also coded, including age, education level (recoded into high, medium, and low), gender, and employment status (binary, recoded into employed or unemployed). The independent variables of interest are (1) national welfare spending as a percentage of GDP, (2) satisfaction with the welfare state, and (3) self-perceived AI replacement risk.

National welfare spending was obtained from the OECD's Social Expenditure (SOCX) database. It covers aggregated social spending information across different programs including unemployment, healthcare, family benefits, and pensions, among others. To make these estimates comparable across country, they are standardized as a percentage of a country's GDP in a given year. I further average them across the ten-year period of 2013-2023.[2] The final variable represents a country-level average of national welfare

[2]Estimates for 2023 are not available for Australia, Canada, New Zealand, Japan, and Costa Rica. For these countries, the variable is taken as the average of spending from 2013-2022.

spending, as a percentage of GDP, for the past decade. Due to a lack of data availability, I was not able to find a disaggregated measure of unemployment benefits specifically.

Subjective perceptions of the welfare state are assessed through the respondent's level of agreement with the following statement: "I think that the government does/would provide my household and me with adequate income support in the case of income loss due to unemployment." The variable is measured on a 5-point scale, and higher values indicate higher levels of agreement.[3] Importantly, this variable is unique to the provision of unemployment benefits, and not the welfare system as a whole.

Self-perceived replacement risk is assessed by an individual's perceived likelihood of their job or job opportunities being taken over by an AI tool like ChatGPT, with responses coded on a 4-point scale from *Very Unlikely* (1) to *Very Likely* (4). The mean response was a 2.1 (*Unlikely*). Interestingly, the occupational sectors with the highest responses on average were mining and quarry, financial and professional services, and water and sewage management. This reinforces the earlier point that many diverse sectors are economically exposed to AI. The full list of occupations, ranked by perceived risk, is available as Figure A.2. in the appendix.

Dependent variables were assessed with six different questions, divided into threat and opportunity perceptions. Respondents were asked to evaluate, on the same 4-point scale, how likely it would be that AI would do each of the following in the foreseeable future:

Threat	Opportunity
• Lead to higher levels of unemployment. • Lead to a rise in income inequality. • Lead to more surveillance at work.	• Create more jobs. • Allow most people to have more free time. • Allow most people to have less boring, repetitive, stressful or mentally demanding jobs.

Descriptive Analyses

I begin with some descriptive analyses to better understand existing attitudes. As shown in Figure 1, every single country sees AI as more of a threat than an opportunity, although the extent to which varies. The average difference between threat and opportunity is about 0.5 units on the 4-point likelihood scale, and the value 3.0 indicates an outcome respondents marked *likely*. These results are mixed, but lean towards pessimism, which is in line with those found in other surveys about attitudes to AI, with a 2025 study by Pew Research Center finding that 34% of adults globally are more concerned than excited, 42% as equally concerned and excited, and 16% as more excited than concerned.

[3]The values 1-5 represent the scale from *Strongly Disagree, Disagree, Neither, Agree, and Strongly Agree.*

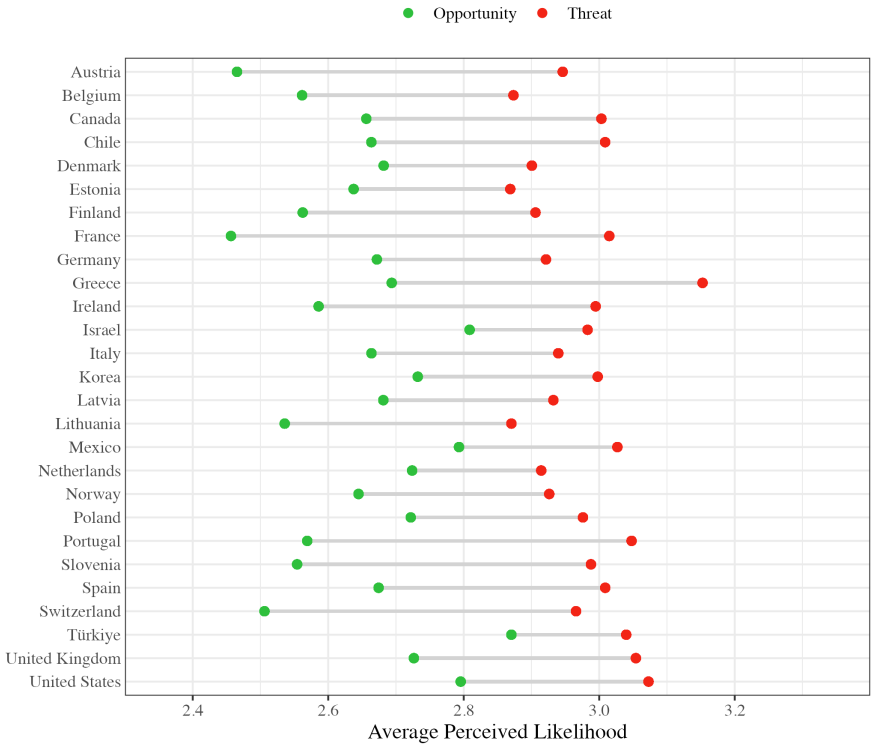

Figure 1: Average perceived likelihood of AI being a threat versus opportunity, by country

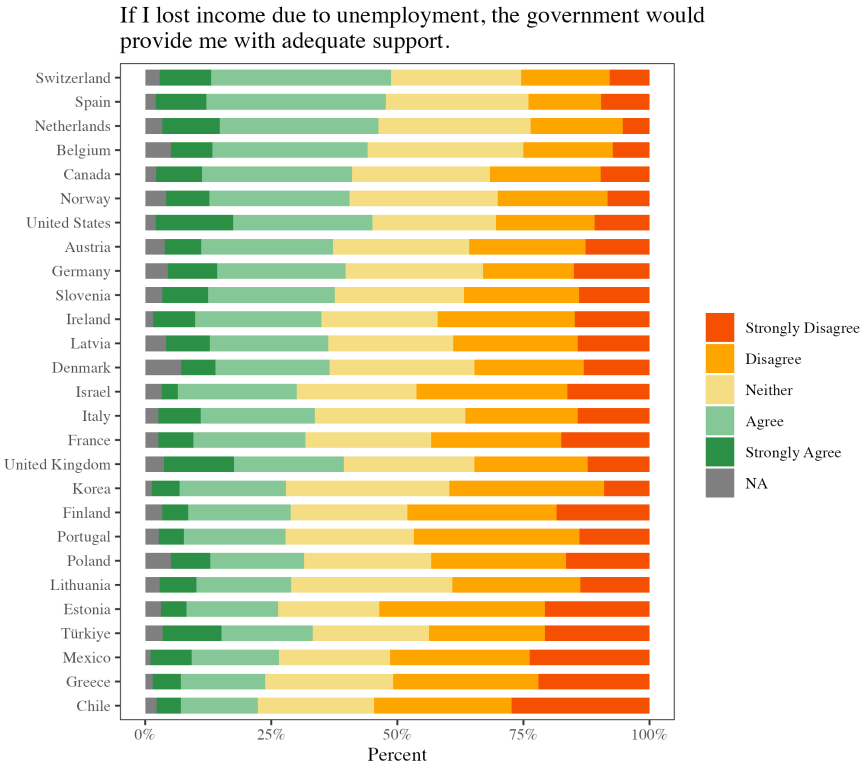

Figure 2: Rate of unemployment support access across OECD countries.

Table 2: Predictors of perceived likelihood of receiving unemployment support.

	Likelihood of Unemployment Benefits
(Intercept)	3.342***
	(0.073)
Age by Decade	−0.155***
	(0.006)
Education Level	0.021
	(0.013)
Male	0.197***
	(0.016)
Income Decile	−0.002
	(0.003)
Num.Obs.	23 240
R2	0.080
R2 Adj.	0.078

+ p < 0.1, * p < 0.05, ** p < 0.01, *** p < 0.001

I then move onto attitudes towards welfare access. When asked whether or not they believe the government would adequately support them if unemployed, I find generally balanced results across all countries, with a slight majority towards disagreement. Figure 2 shows the results, with the countries ranked by decreasing levels of agreement. These roughly cohere with objective levels of welfare generosity, which can be found in Figure A.3. in the Appendix.

Next, I check for predictors of receiving unemployment benefits to exclude the possibility of strong confounders. I find that age and gender are the only variables significantly associated with higher levels of belief. Older individuals are less likely to believe that they would receive unemployment assistance, compared to younger individuals. Men are more likely to believe that they would compared to women. Income and education, somewhat surprisingly, have no significant relationship with this belief. Given the presence of these correlates, my main regression results will control for these demographic predictors.

Model Specifications

Given that my main three independent variables cover both the individual and country level, the data are structured hierarchically. Specifically, individual respondents are clustered within a country and share the same national welfare state, leading the errors to be correlated at the country level. This violates a key assumption of ordinary least squares (OLS) regression. To address this, I use a multilevel model (MLM) with random intercepts at the country-level. This approach reduces bias in the standard errors and improves the precision of my estimates and allows for the inclusion of both national and individual level predictors. In sum, the model I estimate is as follows:

$$
\begin{aligned}
Y_i = \beta_0 &+ \beta_1 \,\text{NationalWelfare}_i + \beta_2 \,\text{LikelihoodBenefits}_i + \beta_3 \,\text{ReplacementRisk}_i \\
&+ \beta_4 \,(\text{NationalWelfare}_i \times \text{ReplacementRisk}_i) \\
&+ \beta_5 \,(\text{LikelihoodBenefits}_i \times \text{ReplacementRisk}_i) \\
&+ \gamma^\top \mathbf{C}_i + u_{c[i]} + \varepsilon_i
\end{aligned}
$$

where C_i is a vector of controls that contains gender, age, education level, employment status, and occupational sector, $u_{c[i]}$ is a random intercept that varies by country, and ε_i is the error term. I run eleven variations of the model above on different outcome variables: six indicating threat/opportunity perceptions, and five for policy preferences.

4. Results

Threat/Opportunity Perceptions

Table 3 shows the predicted marginal effects of being satisfied with the welfare state on the six threat/opportunity outcomes. The results clearly demonstrate that the subjective perception of welfare generosity is much more important than objective levels of welfare spending when individuals evaluate their economic beliefs about AI. To begin, national welfare spending seems to have little association with threat perceptions. Respondents in more generous welfare states do not exhibit statistically significant lower levels of belief in perceptions of AI as increasing unemployment, inequality, or workplace surveillance. However, those in more generous welfare states do have weakly significant lower levels of belief in opportunity. This refutes H1.

However, the evidence for subjective interpretations is much stronger. The results show that the subjective perception of receiving unemployment benefits is strongly associated with both threat and opportunity perceptions in the hypothesized directions ($p < 0.05$). Those who believe that their government will protect them in the case of unemployment are significantly less likely to see AI as a threat across all three outcomes. The highest magnitude relationship is with the belief that AI could increase inequality. They are also more likely to see AI as an opportunity, again at a high level of statistical significance. These results seem to indicate that subjective perceptions of welfare access matter more than objective levels of welfare expenditure. They could also indicate that access to benefits, specifically unemployment benefits, matter independent of the state's generosity.

Turning now to the interaction terms, few are significant for national welfare spending, indicating that the relationship between national welfare expenditure and the outcome variables is the same at different levels of perceived replacement risk. However, the terms are positive and statistically significant for the interaction between receiving benefits and replacement risk. For threat perceptions, it would seem that the more an individual believes they will be replaced by AI, the less reassured they feel by the protectiveness of unemployment benefits. For opportunity perceptions, the positive effect of unemployment benefits on AI optimism is stronger at higher levels of believed replacement risk.

Table 3: Main regression results for threat and opportunity perceptions.

	Threat			Opportunity		
	Increase Unemployment	Increase Inequality	Enable Workplace Surveillance	Create New Jobs	Increase Free Time	Reduce Boring Work
National Welfare Spending as Percentage of GDP	−0.002	0.000	0.004	−0.009*	−0.012**	−0.007*
	(0.003)	(0.003)	(0.003)	(0.004)	(0.004)	(0.003)
Likelihood of Unemployment Benefits	−0.109***	−0.138***	−0.088***	0.054***	0.046***	0.060***
	(0.012)	(0.012)	(0.011)	(0.012)	(0.011)	(0.012)
Self-Perceived Replacement Risk	0.075**	0.055*	0.031	−0.002	0.035	0.028
	(0.027)	(0.026)	(0.025)	(0.027)	(0.026)	(0.026)
National Welfare x Replacement Risk	0.000	0.000	−0.002+	0.001	0.002+	0.001
	(0.001)	(0.001)	(0.001)	(0.001)	(0.001)	(0.001)
Unemployment Benefits x Replacement Risk	0.038***	0.044***	0.039***	0.056***	0.028***	0.023***
	(0.005)	(0.005)	(0.005)	(0.005)	(0.005)	(0.005)
Num.Obs.	16 857	17 757	17 838	17 529	17 862	17 598
R2 Marg.	0.054	0.053	0.030	0.160	0.090	0.070
R2 Cond.	0.060	0.059	0.042	0.173	0.104	0.077
ICC	0.0	0.0	0.0	0.0	0.0	0.0
Controls	Yes	Yes	Yes	Yes	Yes	Yes

+ p < 0.1, * p < 0.05, ** p < 0.01, *** p < 0.001

The main pattern appears to be that across the board, the likelihood of receiving unemployment benefits reduce threat perceptions and increase opportunity perceptions across the board, while national welfare spending has few significant effects. Feeling personally replaceable by AI increases those threat perceptions, but does not impact opportunity perceptions. I plot predicted marginal effects for levels of perceived unemployment support in Figures 3 and 4 to visually confirm these results. Finally, expected unemployment benefits matter less for those at high replacement risk when the outcome is a threat, and matter more for those at high risk when the outcome is an opportunity. These results are robust to controls for education levels, age, gender, and occupational sector, and country of residence.

Policy Preferences

Table 4 shows the results of my model on policy preferences. The evidence for these hypotheses is much more mixed. As with perceptions, national welfare spending has little substantive relationship with policy preferences of any kind, with most effects close to zero and not statistically significant. However, subjective evaluations of welfare do matter, though not in the direction originally predicted. Those who believe they would receive unemployment benefits are prefer less government regulation almost across the board, with the exception of investing in infrastructure.

The divergence theory appears not to extend to policy preferences. This might point to the fact that those who believe they would receive sufficient benefits feel as though the system is sufficient as is, and have no desire to see increased government intervention. Self-perceived replacement risk has little substantive effect on policy outcomes, with the surprising

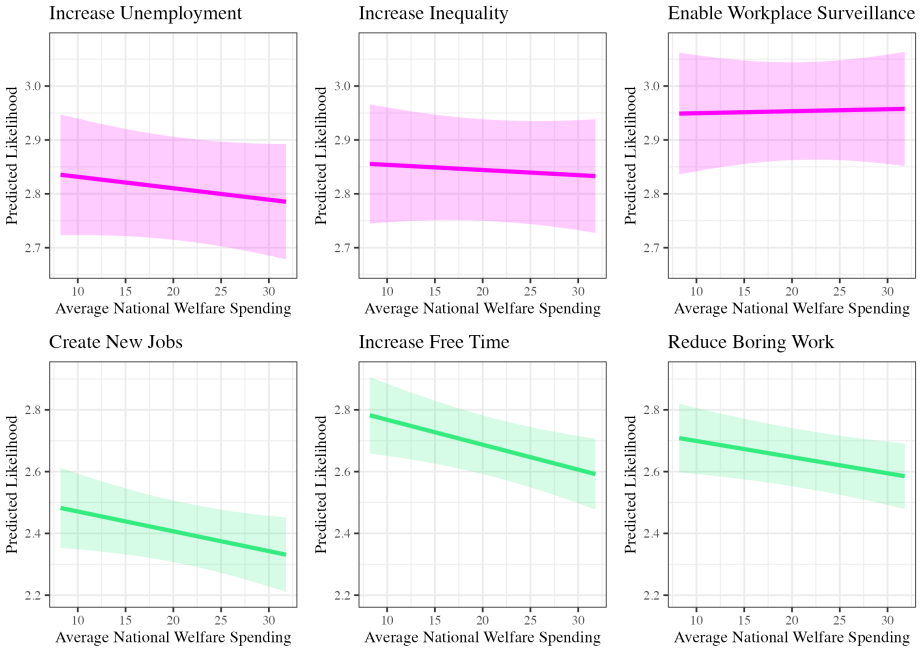

Figure 3: Risk perceptions based on avergage levels of welfare expenditure (as % of GDP).

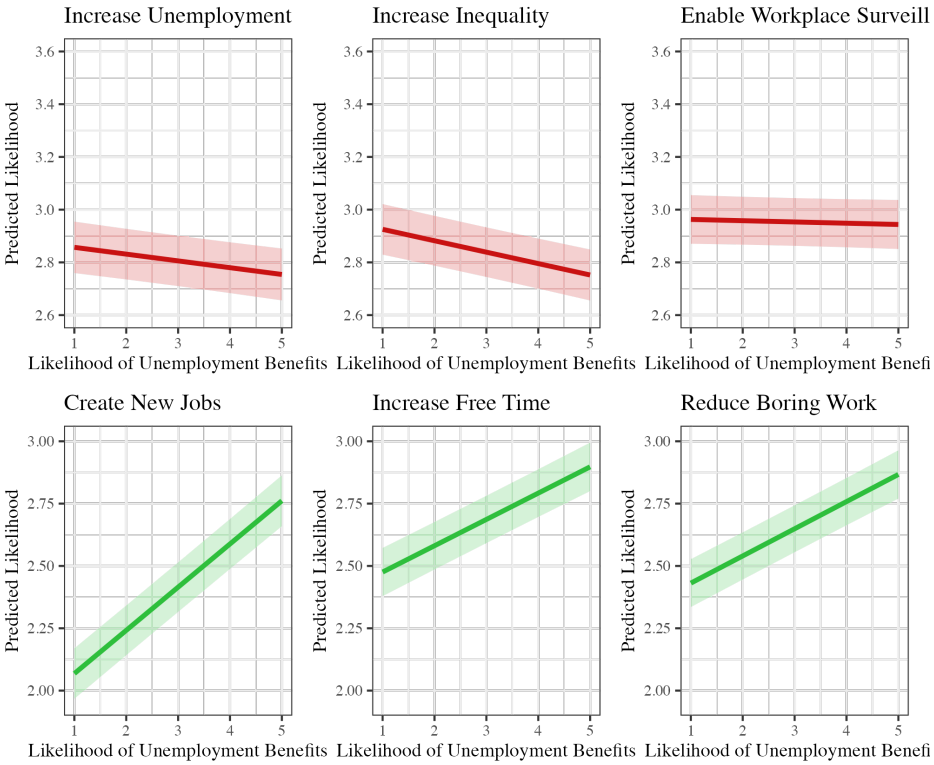

Figure 4: Predicted marginal effects at different levels of receiving unemployment benefits.

Table 4: Regression results for corresponding policy preferences.

	Integration			Regulation	
	Retrain Workers	Invest in Digital Infrastructure	Subsidize Hardest-Hit Companies	Introduce Tariffs	Introduce Robot Tax
National Welfare Spending as Percentage of GDP	0.000	−0.007	−0.006	−0.015*	−0.011+
	(0.004)	(0.005)	(0.005)	(0.006)	(0.006)
Likelihood of Unemployment Benefits	−0.088***	0.001	−0.034*	−0.036*	−0.069***
	(0.012)	(0.013)	(0.014)	(0.014)	(0.015)
Self-Perceived Replacement Risk	−0.089***	−0.026	−0.012	−0.105**	−0.048
	(0.027)	(0.028)	(0.031)	(0.032)	(0.034)
National Welfare x Replacement Risk	−0.002*	−0.001	−0.001	0.002+	0.002
	(0.001)	(0.001)	(0.001)	(0.001)	(0.001)
Unemployment Benefits x Replacement Risk	0.039***	0.023***	0.036***	0.052***	0.051***
	(0.005)	(0.005)	(0.006)	(0.006)	(0.006)
Num.Obs.	18 639	18 269	18 385	18 036	18 110
R2 Marg.	0.023	0.042	0.023	0.033	0.033
R2 Cond.	0.036	0.073	0.038	0.066	0.053
ICC	0.0	0.0	0.0	0.0	0.0
Controls	Yes	Yes	Yes	Yes	Yes

+ $p < 0.1$, * $p < 0.05$, ** $p < 0.01$, *** $p < 0.001$

exception of decreasing support for retraining programs.

However, the positive interaction terms for unemployment benefits and replacement risk are interesting. Perceived unemployment insurance and AI risk each individually reduce support for government intervention of any kind, but together they shift attitudes in the opposite direction. Economic security appears to make AI risk politically actionable: when people feel both threatened and institutionally supported, they are more willing to back proactive government policies.

5. Conclusion

This paper set out to evaluate whether institutional capacities, like welfare state generosity, can pre-emptively shape economic outlooks like those towards AI, in a predicted pattern of threat or opportunity divergence. Across my analysis, the results point to the primacy of subjective, rather than objective, evaluations of welfare as a strong determinant of economic outlooks.

To begin, national welfare generosity played a surprisingly limited role. Residing in a more generous welfare state had little association with an individual's AI threat perceptions, and in fact, a negative association with their opportunity perceptions. Further research will need to investigate why citizens who are objectively afforded a higher level of social protection feel more pessimistic about their economic outlooks with regards to AI.

However, subjective assessments of welfare access played a substantively large role in predicting diverging perceptions of AI, as expected. Individuals who believed they would receive government unemployment support were consistently more optimistic and less threatened. These results were substantive and robust in across models with different oucomes and with full sets of demographic controls. This suggests that the consequences of welfare institutions operate less through national spending levels and more through whether individuals believe those institutions will work for them personally.

Importantly, the interaction between perceived unemployment protection and AI replacement risk exhibits a conditionality. For threat perceptions, unemployment insurance is less reassuring among those who feel highly vulnerable to AI displacement, indicating that perceived risk can outweigh the stabilizing effects of welfare access. For opportunity perceptions, however, welfare expectations and perceived risk positively reinforce one another: individuals who feel both protected and threatened are especially likely to view AI as an opportunity. These patterns suggest that economic security does not simply dampen concern, but rather reshapes how risk is interpreted.

Turning to policy preferences, the results are a different story. Neither national welfare spending nor perceived unemployment protection consistently increases support for government intervention. Instead, individuals who expect to receive unemployment benefits tend to oppose additional intervention of any kind, whether oriented towards regulation or integration. This pattern implies that confidence in existing welfare arrangements may reduce demand for further state action. At the same time, self-perceived replacement risk indepen-

dently reduces support for retraining programs, underscoring that feeling vulnerable does not automatically translate into support for intervention.

Yet here again, the interaction between unemployment benefit expectations and AI risk is important. While each factor alone is associated with lower support for intervention, their combination reverses this effect. When individuals feel both economically protected and personally threatened by AI, support for proactive policies—particularly retraining—rises. Economic security, in this sense, appears to make technological risk politically actionable rather than politically paralyzing, although the direction of political action does not diverge in my hypothesized directions.

Several avenues for future studies are possible. Further research could evaluate the effect of objective levels of risk or exposure through one of many assessment metrics published in recent economics literature, as opposed to self-perceived risk, which some social scientific literature argues is frequently underestimated (Zhang, 2022). Other areas for further analysis could study the surprising negative relationship between welfare state generosity and opportunity perceptions.

In sum, these findings suggest that welfare states matter for economic outlooks towards AI not through their size and generosity, but primarily through citizens' beliefs about access and reliability. Responses to automation are affected less by what welfare states spend, and more by whether or not people believe that the safety net will be there when they need it. This suggests that expanding welfare capacity alone may be insufficient to shape attitudes toward technological change, unless accompanied by institutions and communication that make social protection visible, credible, and personally relevant.

References

Acemoglu, D., & Restrepo, P. (2019). Automation and new tasks: How technology displaces and reinstates labor. Journal of Economic Perspectives, 33(2), 3–30. https://doi.org/10.1257/jep.33.2.3

Antràs, P. (2020). De-Globalisation? Global value chains in the post-covid-19 age. National Bureau of Economic Research. https://doi.org/10.3386/w28115

Autor, D., Dorn, D., Hanson, G., & Majlesi, K. (2020). Importing political polarization? The electoral consequences of rising trade exposure. American Economic Review, 110(10), 3139–3183. https://doi.org/10.1257/aer.20170011

Autor, D. H., Dorn, D., & Hanson, G. H. (2013). The china syndrome: Local labor market effects of import competition in the united states. American Economic Review, 103(6), 2121–2168. https://doi.org/10.1257/aer.103.6.2121

Autor, D. H., Levy, F., & Murnane, R. J. (2003). The skill content of recent technological change: An empirical exploration. The Quarterly Journal of Economics, 118(4), 1279–1333. https://doi.org/10.1162/003355303322552801

Gallego, A., Kurer, T., & Schöll, N. (2022). Neither left behind nor superstar: Ordinary winners of digitalization at the ballot box. The Journal of Politics, 84(1), 418–436. https://doi.org/10.1086/714920

Katz, L. F., & Murphy, K. M. (1992). Changes in relative wages, 1963-1987: Supply and demand factors. The Quarterly Journal of Economics, 107(1), 35–78. https://doi.org/10.2307/2118323

McElheran, K., Li, J. F., Brynjolfsson, E., Kroff, Z., Dinlersoz, E., Foster, L., & Zolas, N. (2024). AI adoption in America: Who, what, and where. Journal of Economics & Management Strategy, 33(2), 375–415. https://doi.org/10.1111/jems.12576

Michaels, G., Natraj, A., & Van Reenen, J. (2014). Has ICT polarized skill demand? Evidence from eleven countries over twenty-five years. The Review of Economics and Statistics, 96(1), 60–77. https://doi.org/10.2307/43554913

OECD. (2024a). OECD Risks that Matter Survey. https://www.oecd.org/en/about/programmes/oecd-risks-that-matter-rtm-survey.html#rtm-data-and-methodology

OECD. (2024b). OECD Social Expenditure (SOCX) Database.

Poushter, J., Fagan, M., & Corichi, M. (2025, October 15). 2. Concern and excitement about AI. Pew Research Center. https://www.pewresearch.org/global/2025/10/15/concern-and-excitement-about-ai/

Rodrik, D. (2021). Why does globalization fuel populism? Economics, culture, and the rise of right-wing populism. Annual Review of Economics, 13(1), 133–170. https://doi.org/10.1146/annurev-economics-070220-032416

Vlandas, T., & Halikiopoulou, D. (2021). Welfare state policies and far right party support: Moderating 'insecurity effects' among different social groups. West European Politics, 45(1), 24–49. https://doi.org/10.1080/01402382.2021.1886498

Webb, M. (2020). The impact of artificial intelligence on the labor market. https://papers.ssrn.com/sol3/papers.cfm?abstract_id=3482150

Appendix

Figure A.1. Self-perceived risk of replacement by occupational sector.

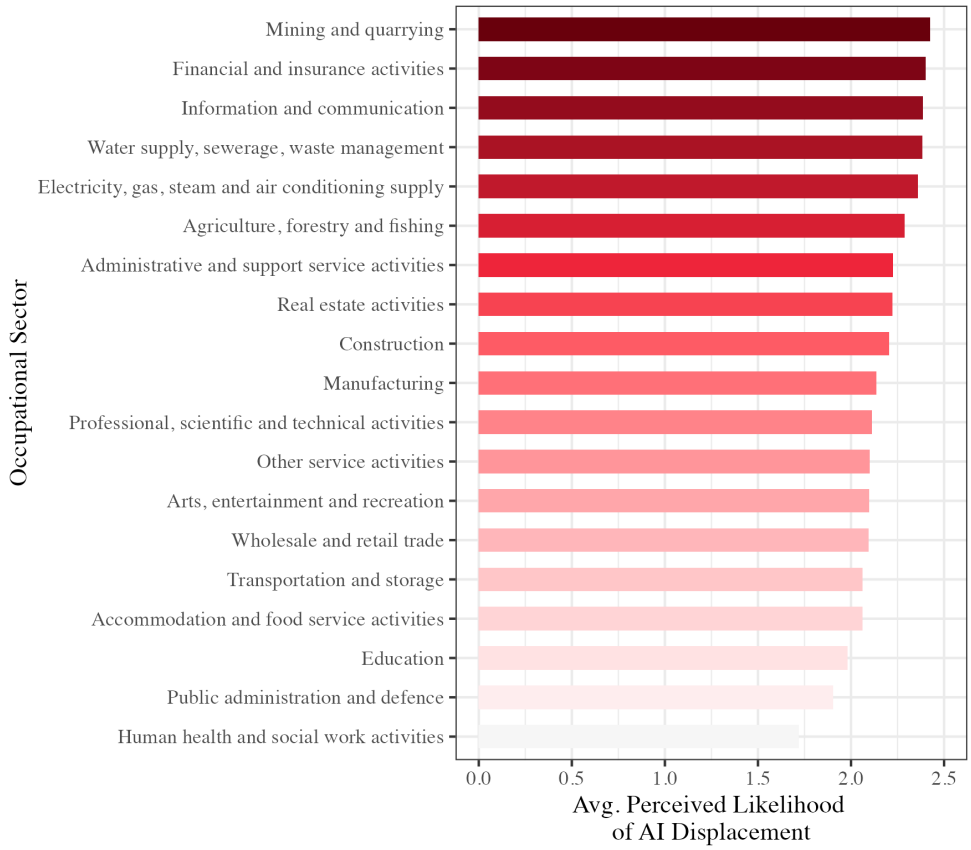

Figure A.2. Summary Statistics for all Variables.

Variable	Mean	SD	Min	Max
Independent Variables				
Age	3.13	1.33	1	5
AI Replacement Risk	2.12	0.99	1	4
Education	2.36	0.66	1	3
Employed	0.80	0.40	0	1
Male	1.50	0.50	1	2
Would Receive Unemployment	2.87	1.19	1	5
Dependent Variables				
AI Increases Free Time	2.77	0.83	1	4
AI Increases Inequality	2.91	0.84	1	4
AI Creates New Jobs	2.44	0.90	1	4
AI Reduces Boring Work	2.71	0.84	1	4
AI Enables Workplace Surveillance	3.05	0.78	1	4
AI Increases Unemployment	2.97	0.84	1	4

Figure A.3.

Table 5: OECD countries ranked by national welfare expenditure as % of GDP, 2013-2023.

Country	Avg Welfare Spending
France	31.8
Finland	30.4
Austria	29.3
Belgium	28.7
Denmark	28.7
Italy	28.4
Sweden	27.0
Germany	26.8
Spain	25.6
Greece	25.6
Portugal	24.0
Norway	23.7
United Kingdom	23.5
Slovenia	23.1
Japan	22.9
New Zealand	21.7
Luxembourg	21.4
Poland	21.1

Country	Avg Welfare Spending
Czechia	20.3
Canada	19.9
Hungary	19.3
United States	19.3
Iceland	18.8
Netherlands	18.7
Australia	18.3
Slovak Republic	17.9
Latvia	17.8
Estonia	17.5
Lithuania	16.9
Israel	16.3
Switzerland	16.2
Ireland	14.9
Colombia	14.4
Chile	12.5
Costa Rica	12.4
Korea	11.6
Türkiye	11.5
Mexico	8.2